



Back to Backprop

Neural networks from scratch

Jovan Krajevski

May 2025

Table of Contents

1 Introduction

► Introduction


- Linear regression
- More predictors
- The matrix form

- Linear projections
- Linear layers are not enough
- Finding the best fit
- The computational graph
- Live coding experience

What even is a neural network?

1 Introduction


Biological neuron



Don't use
this comparison
in your slides.

And **NEVER** use
this arrow!

Artificial neuron



img source: Stop using biological analogies to describe AI. It's 99.999% wrong.

Why the biological analogy?

1 Introduction

- It is supposed to be useful... as useful as:
 - the "car is an artificial horse" analogy
 - the "plane is an artificial bird" analogy
- But it is compelling...
 - the road to artificial "intelligence" is paved with artificial "neurons"
- ...and clouds the judgement when doing research

Neurons as calculators

1 Introduction

- Neurons can multiply numbers
- Neurons can add numbers
- Neurons can choose the larger number
- But they usually can't do a lot more
- Neurons are functions
 - Multiple inputs can be related to the same output
 - Only one output can be related to a given input

The purpose of this lecture

1 Introduction

Boring reasons

- Know what's under the hood as an intellectual curiosity
- Improve on the core algorithm

Practical reasons

- Backprop is a leaky abstraction
- Develop a mathematical intuition useful for research/debugging

- Vanishing gradients on sigmoids (or tanh)
- Dead ReLUs
- Karpathy: Yes you should understand backprop

Table of Contents

2 Linear regression

- ▶ Introduction
- ▶ **Linear regression**
- ▶ More predictors
- ▶ The matrix form
- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ Live coding experience

A linear function


2 Linear regression

$$y = xk + m$$

- If the plane was a grid
 - m is where you start
 - if you move one block to the right, you move k blocks up
- k - the slope - the weight - the rate of change
- m - the intercept - the bias - the intersection with y -axis
 - The intersection occurs when $x = 0$
 - $x = 0 \implies y = 0 \cdot k + m = m$


Tweaking the slope

2 Linear regression



Tweaking the intercept

2 Linear regression




Parameters

2 Linear regression

- The slope and the intercept - parameters
- Every straight line can be expressed by tweaking k and m
 - Except the vertical line; why?
- So why is this useful?

Let's look at some data

2 Linear regression



The task of linear regression


2 Linear regression

- Fit a linear function to the data
 - Find values for k and m that approximate the data
 - Use k and m to make out-of-sample predictions
- What is a good fit?
 - For each sample $(x_i, y_i); i \in \mathbb{N}, i < N$ and fixed values for $k = \hat{k}$ and $m = \hat{m}$ calculate the distance between $\hat{y} = x\hat{k} + \hat{m}$ and y_i
 - $err_i = |y_i - \hat{y}_i|$ or $error_i = (y_i - \hat{y}_i)^2$
 - $err_{avg} = \sum_{i=1}^N err_i / N$
- We need to optimize:

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Some "fits"

2 Linear regression



The errors

2 Linear regression



Table of Contents

3 More predictors

- ▶ Introduction
- ▶ Linear regression
- ▶ **More predictors**
- ▶ The matrix form
- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ Live coding experience

Multiple predictors

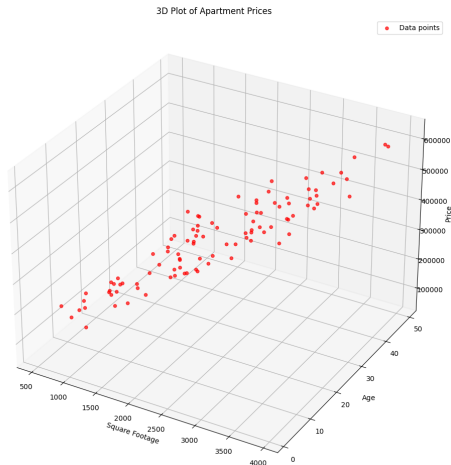
3 More predictors

- x_i - predictor - regressor - attribute - independent variable - feature?
- y_i - target - dependent variable
- We assumed that only the sq footage is available to us
 - But what if we have multiple predictors, like apartment age, floor number, city, location?

$$\mathbf{x}_i = \begin{pmatrix} x_i^{(1)} & x_i^{(2)} & \dots & x_i^{(D)} \end{pmatrix}$$


Let's look at some 3D data

3 More predictors



Let's "fit" that 3D data

3 More predictors



Linear regression with multiple predictors

3 More predictors

- Let's change the notation a little bit
 - let the slope be w - *weight*
 - let the intercept be b - *bias*

$$\hat{y}_i = x_i w + b$$

- For multiple predictors:

$$\hat{y}_i = x_i^{(1)} w_1 + x_i^{(2)} w_2 + \dots + x_i^{(D)} w_D + b$$

Table of Contents

4 The matrix form

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form

- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ Live coding experience

And now let's introduce vectors...

4 The matrix form

- Vectors - quantities that have magnitude and direction
- If we look at a vector as a point (we can't really...)
 - magnitude is the distance from the origin
 - direction is always origin \rightarrow vector
- Vectors are finite sequences of a fixed length
 - so we can represent the sample \mathbf{x}_i as a row vector
 - we can also represent the weights \mathbf{w} as a column vector

$$\mathbf{x}_i = \begin{pmatrix} x_i^{(1)} & x_i^{(2)} & \dots & x_i^{(D)} \end{pmatrix}, \mathbf{w} = \begin{pmatrix} w_1 \\ w_2 \\ \vdots \\ w_D \end{pmatrix}$$

Why bother with vectors?

4 The matrix form

- Because of the vector operations (they are faster)
- Because of the benefits of linear algebra
- The dot (inner) product

$$\begin{aligned}\mathbf{x}_i \mathbf{w} &= x_i^{(1)} w_1 + x_i^{(2)} w_2 + \dots + x_i^{(D)} w_D \\ &\implies \hat{y}_i = \mathbf{x}_i \mathbf{w} + b\end{aligned}$$

Representing data in matrix form

4 The matrix form

- A matrix is a rectangular array; you can think of it as:
 - a row vector consisting of column vectors
 - a column vector of row vectors

$$\mathbf{X} = \begin{pmatrix} \text{---}\mathbf{x}_1\text{---} \\ \text{---}\mathbf{x}_2\text{---} \\ \vdots \\ \text{---}\mathbf{x}_N\text{---} \end{pmatrix} = \begin{pmatrix} \begin{array}{c} | \\ \mathbf{x}^{(1)} \\ | \end{array} & \begin{array}{c} | \\ \mathbf{x}^{(2)} \\ | \end{array} & \dots & \begin{array}{c} | \\ \mathbf{x}^{(D)} \\ | \end{array} \end{pmatrix}$$

- Rows are samples, columns are predictors!
- What if we multiplied $\mathbf{X}\mathbf{w}$?

The matrix-vector product

4 The matrix form

$$\mathbf{X}\mathbf{w} = \begin{pmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \vdots \\ \mathbf{x}_N \end{pmatrix} \begin{pmatrix} w_1 \\ w_2 \\ \vdots \\ w_D \end{pmatrix} = \begin{pmatrix} \mathbf{x}_1\mathbf{w} \\ \mathbf{x}_2\mathbf{w} \\ \vdots \\ \mathbf{x}_N\mathbf{w} \end{pmatrix} = \begin{pmatrix} x_1^{(1)}w_1 + x_1^{(2)}w_2 + \dots + x_1^{(D)}w_D \\ x_2^{(1)}w_1 + x_2^{(2)}w_2 + \dots + x_2^{(D)}w_D \\ \vdots \\ x_N^{(1)}w_1 + x_N^{(2)}w_2 + \dots + x_N^{(D)}w_D \end{pmatrix}$$

$$\hat{\mathbf{y}} = \begin{pmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_N \end{pmatrix} = \mathbf{X}\mathbf{w} + b$$

Table of Contents

5 Linear projections

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form
- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ Live coding experience

Let us look at a different problem now

5 Linear projections

- Forget about "fitting" for a second...
- Let \mathbf{X} be $N \times 3$ matrix representing real estate data
 - column 1 - sq footage
 - column 2 - number of bedrooms
 - column 3 - age
- We are interested in estimating:
 - $\mathbf{h}^{(1)}$ - space and comfort
 - $\mathbf{h}^{(2)}$ - property condition

Multiple linear regressions

5 Linear projections

$$\mathbf{w}_1 = \begin{pmatrix} 0.8 \\ 0.6 \\ -0.2 \end{pmatrix} \quad b^{(1)} = 0.3 \quad \mathbf{w}_2 = \begin{pmatrix} 0.2 \\ 0.1 \\ -1.2 \end{pmatrix} \quad b^{(2)} = -0.7$$

$$\mathbf{h}^{(1)} = \mathbf{X}\mathbf{w}_1 + b^{(1)} \quad \mathbf{h}^{(2)} = \mathbf{X}\mathbf{w}_2 + b^{(2)}$$

Representing weights in matrix form

5 Linear projections

- Let M be the number of linear regressions

$$\mathbf{W} = \begin{pmatrix} \text{---}\mathbf{w}_1\text{---} \\ \text{---}\mathbf{w}_2\text{---} \\ \vdots \\ \text{---}\mathbf{w}_D\text{---} \end{pmatrix} = \begin{pmatrix} \begin{array}{c} | \\ \mathbf{w}^{(1)} \\ | \end{array} & \begin{array}{c} | \\ \mathbf{w}^{(2)} \\ | \end{array} & \dots & \begin{array}{c} | \\ \mathbf{w}^{(M)} \\ | \end{array} \end{pmatrix}$$

- What if we multiplied \mathbf{XW} ?
- And maybe created a row vector $\mathbf{b} = \begin{pmatrix} b^{(1)} & b^{(2)} & \dots & b^{(M)} \end{pmatrix}$?

The matrix product

5 Linear projections

$$\begin{aligned}\mathbf{X}\mathbf{W} &= \begin{pmatrix} \text{---}\mathbf{x}_1\text{---} \\ \text{---}\mathbf{x}_2\text{---} \\ \vdots \\ \text{---}\mathbf{x}_N\text{---} \end{pmatrix} \begin{pmatrix} \left. \begin{array}{c} | \\ \mathbf{w}^{(1)} \\ | \end{array} \right| & \left. \begin{array}{c} | \\ \mathbf{w}^{(2)} \\ | \end{array} \right| & \dots & \left. \begin{array}{c} | \\ \mathbf{w}^{(M)} \\ | \end{array} \right| \end{pmatrix} \\ &= \begin{pmatrix} \mathbf{x}_1\mathbf{w}^{(1)} & \mathbf{x}_1\mathbf{w}^{(2)} & \dots & \mathbf{x}_1\mathbf{w}^{(M)} \\ \mathbf{x}_2\mathbf{w}^{(1)} & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \mathbf{x}_{N-1}\mathbf{w}^{(M)} \\ \mathbf{x}_N\mathbf{w}^{(1)} & \dots & \mathbf{x}_N\mathbf{w}^{(M-1)} & \mathbf{x}_N\mathbf{w}^{(M)} \end{pmatrix}\end{aligned}$$

Broadcasting

5 Linear projections

$$\mathbf{X}\mathbf{W} + \mathbf{b} = \begin{pmatrix} \mathbf{x}_1 \mathbf{w}^{(1)} & \mathbf{x}_1 \mathbf{w}^{(2)} & \dots & \mathbf{x}_1 \mathbf{w}^{(M)} \\ \mathbf{x}_2 \mathbf{w}^{(1)} & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \mathbf{x}_{N-1} \mathbf{w}^{(M)} \\ \mathbf{x}_N \mathbf{w}^{(1)} & \dots & \mathbf{x}_N \mathbf{w}^{(M-1)} & \mathbf{x}_N \mathbf{w}^{(M)} \end{pmatrix} + \begin{pmatrix} b^{(1)} & b^{(2)} & \dots & b^{(M)} \end{pmatrix}$$

$$= \begin{pmatrix} \mathbf{x}_1 \mathbf{w}^{(1)} + b^{(1)} & \mathbf{x}_1 \mathbf{w}^{(2)} + b^{(2)} & \dots & \mathbf{x}_1 \mathbf{w}^{(M)} + b^{(M)} \\ \mathbf{x}_2 \mathbf{w}^{(1)} + b^{(1)} & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \mathbf{x}_{N-1} \mathbf{w}^{(M)} + b^{(M)} \\ \mathbf{x}_N \mathbf{w}^{(1)} + b^{(1)} & \dots & \mathbf{x}_N \mathbf{w}^{(M-1)} + b^{(M-1)} & \mathbf{x}_N \mathbf{w}^{(M)} + b^{(M)} \end{pmatrix}$$

Linear projection

5 Linear projections

- $\mathbf{XW} + \mathbf{b}$ - multiple linear regressions - linear projection
 - a.k.a. a linear layer
- It projects data in a new space
- Useful for:
 - Feature extraction
 - Linear separability
 - Data compression (if $M < D$) or expansion (if $M > D$)
- We can "stack" multiple linear projections one after another

”Stacking” linear layers

5 Linear projections

- Let L be the number of linear layers
 - space dims: $M_0 = D, M_1, M_2, \dots, M_L = 1$
 - weights: $\mathbf{W}_1, \mathbf{W}_2, \dots, \mathbf{W}_L$; \mathbf{W}_i is a $M_{i-1} \times M_i$ matrix
 - biases: $\mathbf{b}_1, \mathbf{b}_2, \dots, \mathbf{b}_L$; \mathbf{b}_i is a M_i dimensional row vector
 - $N \times M_i$ matrix after broadcasting!!!
 - outputs: $\mathbf{H}_1, \mathbf{H}_2, \dots, \mathbf{H}_L = \hat{\mathbf{y}}$; \mathbf{H}_i is a $N \times M_i$ matrix

$$\mathbf{H}_1 = \mathbf{X}\mathbf{W}_1 + \mathbf{b}_1$$


$$\mathbf{H}_2 = \mathbf{H}_1\mathbf{W}_2 + \mathbf{b}_2$$

$$\vdots$$

$$\hat{\mathbf{y}} = \mathbf{H}_L = \mathbf{H}_{L-1}\mathbf{W}_L + \mathbf{b}_L$$

Let's visualize this

5 Linear projections




Input Layer $\in \mathbb{R}^2$

Output Layer $\in \mathbb{R}^1$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections




Input Layer $\in \mathbb{R}^2$

Output Layer $\in \mathbb{R}^2$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections




Input Layer $\in \mathbb{R}^2$

Output Layer $\in \mathbb{R}^3$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections




Input Layer $\in \mathbb{R}^2$

Output Layer $\in \mathbb{R}^4$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections



Input Layer $\in \mathbb{R}^2$


Hidden Layer $\in \mathbb{R}^4$

Output Layer $\in \mathbb{R}^1$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections



Input Layer $\in \mathbb{R}^2$


Hidden Layer $\in \mathbb{R}^4$

Output Layer $\in \mathbb{R}^2$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections



Input Layer $\in \mathbb{R}^2$


Hidden Layer $\in \mathbb{R}^4$

Output Layer $\in \mathbb{R}^3$

created with: <https://alexlenail.me/NN-SVG/>

Let's visualize this

5 Linear projections



Input Layer $\in \mathbb{R}^2$

Hidden Layer $\in \mathbb{R}^4$

Hidden Layer $\in \mathbb{R}^3$

Output Layer $\in \mathbb{R}^1$

created with: <https://alexlenail.me/NN-SVG/>

Table of Contents

6 Linear layers are not enough

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form
- ▶ Linear projections
- ▶ **Linear layers are not enough**
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ Live coding experience

Some properties of the matrix product

6 Linear layers are not enough

- Non-commutative: $\mathbf{AB} \neq \mathbf{BA}$
- Associative: $(\mathbf{AB})\mathbf{C} = \mathbf{A}(\mathbf{BC})$
- Distributive: $(\mathbf{A} + \mathbf{B})\mathbf{C} = \mathbf{AC} + \mathbf{BC}$

So what if we stack multiple linear layers?

6 Linear layers are not enough

$$\mathbf{H}_1 = \mathbf{X}\mathbf{W}_1^{[D \times M_1]} + \mathbf{b}_1^{[N \times M_1]}$$

$$\mathbf{H}_2 = \mathbf{H}_1\mathbf{W}_2^{[M_1 \times M_2]} + \mathbf{b}_2^{[N \times M_2]}$$

$$= (\mathbf{X}\mathbf{W}_1^{[D \times M_1]} + \mathbf{b}_1^{[N \times M_1]})\mathbf{W}_2^{[M_1 \times M_2]} + \mathbf{b}_2^{[N \times M_2]} \leftarrow \text{distributive rule}$$

$$= \mathbf{X}\mathbf{W}_1^{[D \times M_1]}\mathbf{W}_2^{[M_1 \times M_2]} + \mathbf{b}_1^{[N \times M_1]}\mathbf{W}_2^{[M_1 \times M_2]} + \mathbf{b}_2^{[N \times M_2]} \leftarrow \text{associative rule}$$

$$= \mathbf{X}\mathbf{Q}_2^{[D \times M_2]} + \mathbf{U}_2^{[N \times M_2]} \leftarrow \text{linear projection!!!}$$

$$\mathbf{H}_i = \mathbf{H}_{i-1}\mathbf{W}_i^{[M_{i-1} \times M_i]} + \mathbf{b}_i^{[N \times M_i]}$$

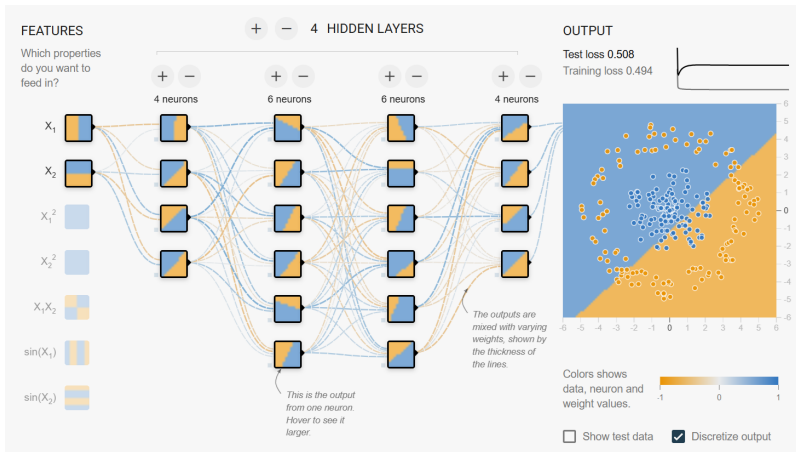
$$= (\mathbf{X}\mathbf{Q}_{i-1}^{[D \times M_{i-1}]} + \mathbf{U}_{i-1}^{[N \times M_{i-1}]})\mathbf{W}_i^{[M_{i-1} \times M_i]} + \mathbf{b}_i^{[N \times M_i]}$$

$$= \mathbf{X}\mathbf{Q}_{i-1}^{[D \times M_{i-1}]}\mathbf{W}_i^{[M_{i-1} \times M_i]} + \mathbf{U}_{i-1}^{[N \times M_{i-1}]}\mathbf{W}_i^{[M_{i-1} \times M_i]} + \mathbf{b}_i^{[N \times M_i]}$$

$$= \mathbf{X}\mathbf{Q}_i^{[D \times M_i]} + \mathbf{U}_i^{[N \times M_i]} \leftarrow \text{linear projection!!!}$$

Stacking multiple linear layers is useless

6 Linear layers are not enough



Non-linearities

6 Linear layers are not enough

- When you stack multiple linear layers, you end up having a linear projection
 - L layers with dims $M_1, \dots, M_L \iff$ 1 layer with dim M_L
 - proof by mathematical induction
- Solution: introduce a non-linear function f between layers - activation
 - can vary depending after which layer it is introduced

$$\mathbf{H}_1 = f(\mathbf{X}\mathbf{W}_1 + \mathbf{b}_1)$$

$$\mathbf{H}_2 = f(\mathbf{H}_1\mathbf{W}_2 + \mathbf{b}_2)$$

$$\vdots$$

$$\hat{\mathbf{y}} = \mathbf{H}_L = f(\mathbf{H}_{L-1}\mathbf{W}_L + \mathbf{b}_L)$$

Elem-wise operations

6 Linear layers are not enough

$$f\left(\begin{pmatrix} \mathbf{x}_1 \mathbf{w}^{(1)} + b^{(1)} & \mathbf{x}_1 \mathbf{w}^{(2)} + b^{(2)} & \dots & \mathbf{x}_1 \mathbf{w}^{(M)} + b^{(M)} \\ \mathbf{x}_2 \mathbf{w}^{(1)} + b^{(1)} & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \mathbf{x}_{N-1} \mathbf{w}^{(M)} + b^{(M)} \\ \mathbf{x}_N \mathbf{w}^{(1)} + b^{(1)} & \dots & \mathbf{x}_N \mathbf{w}^{(M-1)} + b^{(M-1)} & \mathbf{x}_N \mathbf{w}^{(M)} + b^{(M)} \end{pmatrix}\right) =$$
$$\begin{pmatrix} f(\mathbf{x}_1 \mathbf{w}^{(1)} + b^{(1)}) & f(\mathbf{x}_1 \mathbf{w}^{(2)} + b^{(2)}) & \dots & f(\mathbf{x}_1 \mathbf{w}^{(M)} + b^{(M)}) \\ f(\mathbf{x}_2 \mathbf{w}^{(1)} + b^{(1)}) & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & f(\mathbf{x}_{N-1} \mathbf{w}^{(M)} + b^{(M)}) \\ f(\mathbf{x}_N \mathbf{w}^{(1)} + b^{(1)}) & \dots & f(\mathbf{x}_N \mathbf{w}^{(M-1)} + b^{(M-1)}) & f(\mathbf{x}_N \mathbf{w}^{(M)} + b^{(M)}) \end{pmatrix}$$


ReLU

6 Linear layers are not enough

- Popular activations:

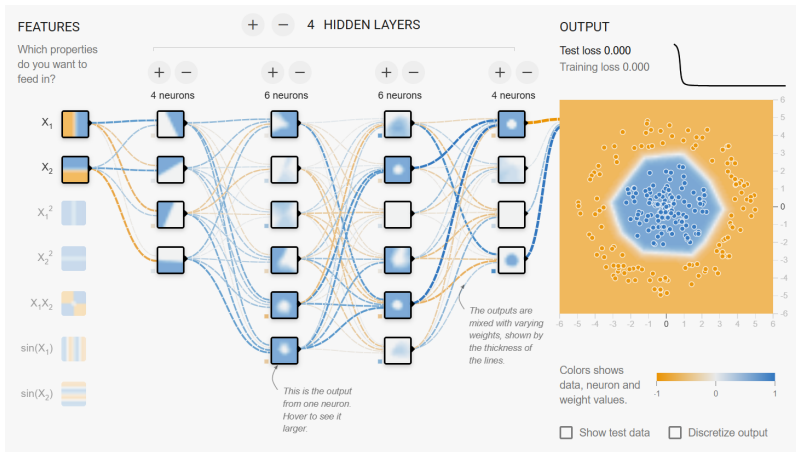
- ReLU
- tanh
- sigmoid

$$\text{ReLU}(x) = \max(0, x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{otherwise} \end{cases}$$



Linear layers + ReLU is useful

6 Linear layers are not enough



created with: Tensorflow Playground

Table of Contents

7 Finding the best fit

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form
- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ **Finding the best fit**
- ▶ The computational graph
- ▶ Live coding experience

The loss function

7 Finding the best fit

$$\begin{aligned}\hat{\mathbf{y}} &= \mathcal{NN}(\mathbf{X}; \mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L) \\ &= \text{ReLU}(\dots \text{ReLU}(\text{ReLU}(\mathbf{X}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2) \dots)\mathbf{W}_L + \mathbf{b}_L\end{aligned}$$

- Parameters: $\theta = \{\mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L\}$
- True values: \mathbf{y} ; Predicted values: $\hat{\mathbf{y}}$
- Find θ such that $\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$ is minimized

Gradient Descent

7 Finding the best fit

Algorithm

1. Choose a random value for $\theta = \hat{\theta}$
2. Calculate $\mathcal{L}(\hat{\theta})$
3. Nudge $\hat{\theta}$ a little bit
4. Repeat from 2

Intuition

1. You are on a field
2. Estimate how low you are
3. Move in a downward direction
4. Repeat from 2


But which way is "downward"?

7 Finding the best fit

- Introducing the derivative - $\frac{d\mathcal{L}}{d\theta}$
 - the rate of change of \mathcal{L} with respect to θ
 - if I change θ a little bit, how much does \mathcal{L} changes?
 - the slope of the tangent of \mathcal{L} in point θ
- Introducing the partial derivatives - $\frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_L}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_L}$
 - if all other parameters are kept the same, what is the rate of change of \mathcal{L} with respect a single parameter?


Let's visualize the derivative

7 Finding the best fit



Let's visualize the derivative

7 Finding the best fit



How much do we "nudge" $\hat{\theta}$?

7 Finding the best fit

- Learning rate - η
 - experimentally determined
 - if η is too large - we skip over the optimum
 - if η is too small - we "fit" too slow

Gradient descent

1. Choose a random value for $\theta = \hat{\theta} = \{\widehat{\mathbf{W}}_1, \dots, \widehat{\mathbf{W}}_L, \widehat{\mathbf{b}}_1, \dots, \widehat{\mathbf{b}}_L\}$
2. Calculate loss $\mathcal{L}(\{\widehat{\mathbf{W}}_1, \dots, \widehat{\mathbf{W}}_L, \widehat{\mathbf{b}}_1, \dots, \widehat{\mathbf{b}}_L\})$ on **complete** dataset \leftarrow *forward* pass
3. Calculate partial derivatives $\frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_L}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_L} \leftarrow$ *backward* pass
4. Update parameters:
$$\widehat{\mathbf{W}}_1 \leftarrow \widehat{\mathbf{W}}_1 - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \dots, \widehat{\mathbf{W}}_L \leftarrow \widehat{\mathbf{W}}_L - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_L}, \widehat{\mathbf{b}}_1 \leftarrow \widehat{\mathbf{b}}_1 - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \dots, \widehat{\mathbf{b}}_L \leftarrow \widehat{\mathbf{b}}_L - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{b}_L}$$
5. Repeat from 2

But nobody really uses gradient descent...

7 Finding the best fit

- Gradient descent is slow - it calculates the loss on the complete dataset before doing an update

Stochastic gradient descent

1. Choose a random value for $\theta = \hat{\theta} = \{\hat{\mathbf{W}}_1, \dots, \hat{\mathbf{W}}_L, \hat{\mathbf{b}}_1, \dots, \hat{\mathbf{b}}_L\}$
2. Choose a random subset of the dataset \leftarrow *batch*
3. Calculate loss $\mathcal{L}(\{\hat{\mathbf{W}}_1, \dots, \hat{\mathbf{W}}_L, \hat{\mathbf{b}}_1, \dots, \hat{\mathbf{b}}_L\})$ on *batch* \leftarrow *forward* pass
4. Calculate partial derivatives $\frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_L}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \dots, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_L} \leftarrow$ *backward* pass
5. Update parameters:
$$\hat{\mathbf{W}}_1 \leftarrow \hat{\mathbf{W}}_1 - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \dots, \hat{\mathbf{W}}_L \leftarrow \hat{\mathbf{W}}_L - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{W}_L}, \hat{\mathbf{b}}_1 \leftarrow \hat{\mathbf{b}}_1 - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \dots, \hat{\mathbf{b}}_L \leftarrow \hat{\mathbf{b}}_L - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{b}_L}$$
6. Repeat from 2

Table of Contents

8 The computational graph

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form
- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ **The computational graph**
- ▶ Live coding experience

Backpropagation in a nutshell

8 The computational graph

- Backpropagation is an efficient way to do a *backward* pass
- Backpropagation = computational graph + *the chain rule*

Transposed vector

8 The computational graph

$$\mathbf{err} = \mathbf{y} - \hat{\mathbf{y}} = \begin{pmatrix} y_1 - \hat{y}_1 \\ y_2 - \hat{y}_2 \\ \vdots \\ y_N - \hat{y}_N \end{pmatrix} \quad (1)$$

- Transposing converts a column vector into a row vector and vice versa
- Transposing "rotates"/switches indices in matrix \mathbf{A} - $a_i^{(j)} \leftarrow a_j^{(i)}$

$$\mathbf{err}^T = (y_1 - \hat{y}_1 \quad y_2 - \hat{y}_2 \quad \dots \quad y_N - \hat{y}_N)$$

Vectorized loss function

8 The computational graph

$$\mathbf{err}^T \mathbf{err} = \begin{pmatrix} y_1 - \hat{y}_1 & y_2 - \hat{y}_2 & \dots & y_N - \hat{y}_N \end{pmatrix} \begin{pmatrix} y_1 - \hat{y}_1 \\ y_2 - \hat{y}_2 \\ \vdots \\ y_N - \hat{y}_N \end{pmatrix} = \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$$\mathcal{L}(\theta) = \frac{\mathbf{err}^T \mathbf{err}}{N} = \frac{(\mathbf{y} - \hat{\mathbf{y}})^T (\mathbf{y} - \hat{\mathbf{y}})}{N}$$

Let's write operations as functions

8 The computational graph

- Let's use a 2 layer NN as an example

$$\hat{\mathbf{y}} = \mathcal{NN}(\mathbf{X}; \mathbf{W}_1, \mathbf{W}_2, \mathbf{b}_1, \mathbf{b}_2) = \text{ReLU}(\mathbf{X}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2$$

- Introduce functions $Add(x, y)$, $Sub(x, y)$, $Mul(x, y)$, $ReLU(x)$, $T(x)$

$$\hat{\mathbf{y}} = Add(Mul(ReLU(Add(Mul(\mathbf{X}, \mathbf{W}_1), \mathbf{b}_1))\mathbf{W}_2), \mathbf{b}_2)$$

$$\mathcal{L}(\theta) = \frac{Mul(T(Sub(\mathbf{y}, \hat{\mathbf{y}})), Sub(\mathbf{y}, \hat{\mathbf{y}}))}{N}$$




So what is a computational graph?

8 The computational graph

- Each node represents a function call
- Each directed edge connects an output of a function/variable to an input of a different function/result

Let's visualize this

8 The computational graph



So what is the chain rule?

8 The computational graph

$$u = g(x) \quad y = f(u) = f(g(x))$$

- Used for compositions of functions
- If we increase u for $\Delta u = 1$, then y changes $\approx \frac{dy}{du}$
- If we increase x for $\Delta x = 1$, then u changes $\approx \frac{du}{dx}$

$$\Delta x = 1 \implies \Delta u \approx \frac{du}{dx} \times 1 \implies \Delta y \approx \frac{dy}{du} \Delta u \approx \frac{dy}{du} \frac{du}{dx} \times 1$$

$$\Delta x \rightarrow 0 \implies \frac{dy}{dx} = \frac{dy}{du} \frac{du}{dx}$$

Chain rule applied to the NN

8 The computational graph

$$\begin{aligned}\mathcal{L} = Mul(u, \dots) &\implies & \frac{\partial \mathcal{L}}{\partial \mathcal{L}} &= 1 \\ u = T(h) &\implies & \frac{\partial \mathcal{L}}{\partial u} &= \frac{\partial \mathcal{L}}{\partial \mathcal{L}} \frac{\partial \mathcal{L}}{\partial u} \\ h = Sub(\dots, \hat{\mathbf{y}}) &\implies & \frac{\partial \mathcal{L}}{\partial h} &= \frac{\partial \mathcal{L}}{\partial u} \frac{\partial \mathcal{L}}{\partial h} \\ & & \frac{\partial \mathcal{L}}{\partial \hat{\mathbf{y}}} &= \frac{\partial \mathcal{L}}{\partial h} \frac{\partial \mathcal{L}}{\partial \hat{\mathbf{y}}}\end{aligned}$$

- And so on...

But that is not the complete chain rule...

8 The computational graph

What if a node is an input to multiple nodes?

$x \rightarrow \text{scalar}$

$$u^{(1)} = g_1(x) \quad u^{(2)} = g_1(x) \quad \dots \quad u^{(n)} = g_n(x)$$

$$u^{(i)} = g_i(x) \quad \mathbf{u} = \begin{pmatrix} u^{(1)} & \dots & u^{(n)} \end{pmatrix}$$

$$y = f(g_1(x), \dots, g_n(x)) = f(\mathbf{u}) \rightarrow \text{not elem-wise!}$$

$y \rightarrow \text{scalar}$

$$\frac{\partial y}{\partial x} = \frac{\partial y}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial x}$$

What are the shapes of the derivatives?

8 The computational graph

- y and x are scalars $\implies \frac{\partial y}{\partial x}$ is scalar
- \mathbf{u} is a column vector $\implies \frac{\partial y}{\partial \mathbf{u}}$ is a row vector - $\left(\frac{\partial y}{\partial u_1} \quad \dots \quad \frac{\partial y}{\partial u_n} \right)$
- \mathbf{u} is a column vector $\implies \frac{\partial \mathbf{u}}{\partial x}$ is a column vector - $\left(\frac{\partial u_1}{\partial x} \quad \dots \quad \frac{\partial u_n}{\partial x} \right)^T$

$$\frac{\partial y}{\partial x} = \left(\frac{\partial y}{\partial u_1} \quad \dots \quad \frac{\partial y}{\partial u_n} \right) \begin{pmatrix} \frac{\partial u_1}{\partial x} \\ \vdots \\ \frac{\partial u_n}{\partial x} \end{pmatrix} = \sum_{i=1}^n \frac{\partial y}{\partial u_i} \frac{\partial u_i}{\partial x}$$

- **Chain rule:** multiply composition, sum if node is input to multiple nodes!

Chain rule in the computational graph

8 The computational graph

- Start at the final node - the gradient is 1
- Pass that gradient to the input nodes - *upstream gradient*
- For each input node:
 - **Chain rule (sum)**: add upstream gradient to *node gradient*
 - Because the node can be an input to multiple nodes!
 - Calculate the gradient of the output with respect to node - *local gradient*
 - **Chain rule (multiply)**: new upstream = upstream \times local
 - Pass new upstream to the input nodes
 - Repeat recursively



Shapes of the gradients

8 The computational graph

- *node gradient* - shape is equal to the node value shape!
- *upstream gradient* - shape is equal to the node value shape!
- *new upstream gradient* - shape is equal to the input node value shape!
- Always check the gradient shapes!

The adding & subtracting derivative

8 The computational graph

- Scalar *Add* first (same for *Sub*, with one minus sign):

$$c = \text{Add}(a, b)$$


$$\Delta a \rightarrow 0 \implies \frac{\partial c}{\partial a} = 1$$

$$\Delta b \rightarrow 0 \implies \frac{\partial c}{\partial b} = 1$$

$$\mathbf{C} = \text{Add}(\mathbf{A}, \mathbf{B}); \mathbf{E} = \text{Sub}(\mathbf{C}, \mathbf{D})$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{C}} = \frac{\partial \mathcal{L}}{\partial \mathbf{E}}; \frac{\partial \mathcal{L}}{\partial \mathbf{D}} = -\frac{\partial \mathcal{L}}{\partial \mathbf{E}}$$

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \frac{\partial \mathcal{L}}{\partial \mathbf{C}}; \frac{\partial \mathcal{L}}{\partial \mathbf{B}} = \frac{\partial \mathcal{L}}{\partial \mathbf{C}}$$



The broadcasting issue

8 The computational graph

- But when we add the biases, the M_i dimensional row vector is broadcasted to a $N \times M_i$ matrix
- The same row vector is added to multiple rows in the $\mathbf{H}_i \mathbf{W}_i$ product
 - It is an input to multiple "nodes"!
- **Chain rule (sum)**: tweaking the biases thus has N times the effect on the loss function
 - node gradient = sum of upstream gradients

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}_i} = \sum_{j=1}^N \frac{\partial \mathcal{L}}{\partial \mathbf{up}_j}$$


Handling multiple outputs

8 The computational graph

- *Sub* has two identical outputs
 - side-effect of graph optimization

$$\mathbf{C} = \text{Sub}(\mathbf{A}, \mathbf{B}); \mathbf{D} = \text{Sub}(\mathbf{A}, \mathbf{B})$$

- Both $\frac{\partial \mathcal{L}}{\partial \mathbf{C}}$ and $\frac{\partial \mathcal{L}}{\partial \mathbf{D}}$ are backpropagated
- **Chain rule (sum):** the final gradient w.r.t. *Sub* is $\frac{\partial \mathcal{L}}{\partial \mathbf{C}} + \frac{\partial \mathcal{L}}{\partial \mathbf{D}}$
- This is equivalent to the broadcasting issue!




The matrix product derivative

- Scalars first:

$$c = Mul(a, b)$$

$$\Delta b \rightarrow 0 \implies \frac{\partial c}{\partial b} = a$$



The chain rule trick

8 The computational graph

\mathcal{L} is a scalar; $\mathbf{C}^{[n \times m]} = \text{Mul}(\mathbf{A}^{[n \times p]}, \mathbf{B}^{[p \times m]}); \quad \mathbf{UP}^{[n \times m]}$

$$\mathbf{c}_i = \begin{pmatrix} a_i^{(1)} b_1^{(1)} & a_i^{(1)} b_1^{(i)} & a_i^{(1)} b_1^{(m)} \\ + & + & + \\ \vdots & \vdots & \vdots \\ + & + & + \\ a_i^{(j)} \mathbf{b}_j^{(1)} & \dots & a_i^{(j)} \mathbf{b}_j^{(i)} & \dots & a_i^{(j)} \mathbf{b}_j^{(m)} \\ + & + & + \\ \vdots & \vdots & \vdots \\ + & + & + \\ a_i^{(p)} b_p^{(1)} & \dots & a_i^{(p)} b_p^{(i)} & \dots & a_i^{(p)} b_p^{(m)} \end{pmatrix}$$

Changing $a_i^{(j)}$ affects row i in \mathbf{C}

$$c_i^{(j)} = \sum_{k=1}^p a_i^{(k)} b_k^{(i)} = \dots + a_i^{(j)} b_j^{(i)} + \dots$$

$$\implies b_j^{(i)} \text{ is the derivative!}$$

The chain rule trick cont'd.

8 The computational graph

- $a_i^{(j)}$ affect all columns in \mathbf{c}_i - input to multiple "nodes" - **chain rule (sum)**!
- Each column in \mathbf{c}_i has upstream gradient - **chain rule (multiply)**!

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial a_i^{(j)}} &= \sum_{k=1}^m \frac{\partial \mathcal{L}}{\partial c_i^{(k)}} \frac{\partial c_i^{(k)}}{\partial a_i^{(j)}} = \sum_{k=1}^m \frac{\partial \mathcal{L}}{\partial c_i^{(k)}} b_j^{(k)} = \frac{\partial \mathcal{L}}{\partial \mathbf{c}_i} \mathbf{b}_j^T \\ &\implies \frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \frac{\partial \mathcal{L}}{\partial \mathbf{C}} \mathbf{B}^T\end{aligned}$$


$$\frac{\partial \mathcal{L}}{\partial \mathbf{B}} = \mathbf{A}^T \frac{\partial \mathcal{L}}{\partial \mathbf{C}} \rightarrow \text{exercise for the reader!}$$

The transpose derivative

8 The computational graph

- Transposing does not "change" the input - it rearranges it
 - $\mathbf{B} = T(\mathbf{A})$
 - $b_i^{(j)} = a_j^{(i)}$
- The derivative shows how the input was rearranged

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \left(\frac{\partial \mathcal{L}}{\partial \mathbf{B}} \right)^T$$




The *ReLU* derivative

8 The computational graph

$$\mathbf{B} = \text{ReLU}(\mathbf{A})$$

- When $\mathbf{A} \leq 0$, the rate of change is 0
 - *ReLU* is flat
- When $\mathbf{A} > 0$, the rate of change is 1
 - *ReLU* is a linear function with slope 1
- \odot is Hadamard product - element-wise product of two matrices

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}} = \frac{\partial \mathcal{L}}{\partial \mathbf{B}} \odot \mathbf{R} \text{ where } r_i^{(j)} = \begin{cases} 0 & \text{if } a_i^{(j)} \leq 0 \\ 1 & \text{if } a_i^{(j)} > 0 \end{cases}$$



The complete backprop graph

8 The computational graph

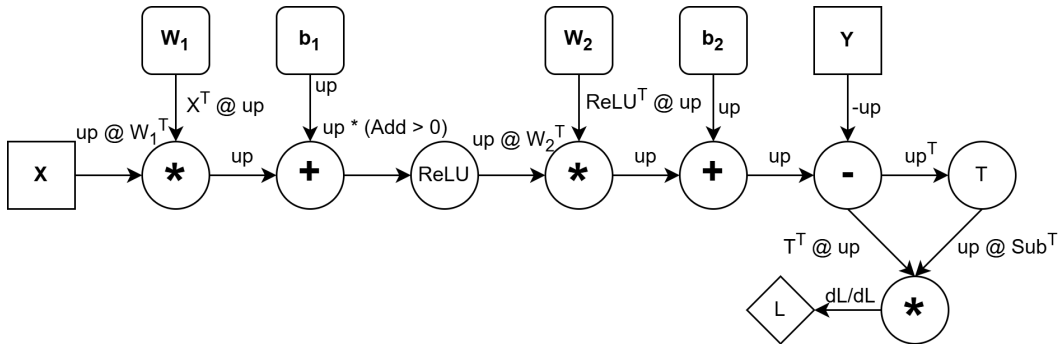


Table of Contents

9 Live coding experience

- ▶ Introduction
- ▶ Linear regression
- ▶ More predictors
- ▶ The matrix form

- ▶ Linear projections
- ▶ Linear layers are not enough
- ▶ Finding the best fit
- ▶ The computational graph
- ▶ **Live coding experience**

Q&A

Thank you for listening!
Your feedback will be highly appreciated!